

Summary

- Built an end-to-end ML project on AI impact on jobs (2030).
- Cleaned and prepared raw data using pandas and NumPy.
- Performed EDA with charts and correlation analysis.
- Trained 5 supervised ML models for job-risk prediction.
- Compared models using accuracy and confusion matrices.
- Applied K-Means clustering for unsupervised job grouping.
- Used PCA for 3D cluster visualization.
- Analyzed feature importance with Random Forest.

- Import all necessary libraries for data analysis and modeling.....

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, learning_curve
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.decomposition import PCA
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
plt.style.use("seaborn-v0_8")
K_VAL = 3

from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
import warnings
warnings.filterwarnings("ignore")
```

- Did Some Exploratory Data Analysis(EDA)....
- Rename columns for consistency.....

```
In [3]: df.columns = [c.replace(" ", "_") for c in df.columns]
```

- Reading dataset using pandas.....

```
In [2]: df = pd.read_csv("D:\Download\AI_Impact_on_Jobs_2030.csv")
```

- Did Some Exploratory Data Analysis....

- Dataset Shape....

```
In [3]: print('Shape:', df.shape)
```

Shape: (3000, 18)

- Dataset Preview

```
In [44]: print(df.head())
```

	Job_Title	Average_Salary	Years_Experience	Education_Level	\
0	Security Guard	45795	28	Master's	
1	Research Scientist	133355	20	PhD	
2	Construction Worker	146216	2	High School	
3	Software Engineer	136530	13	PhD	
4	Financial Analyst	70397	22	High School	

	AI_Exposure_Index	Tech_Growth_Factor	Automation_Probability_2030	\
0	0.18	1.28	0.85	
1	0.62	1.11	0.05	
2	0.86	1.18	0.81	
3	0.39	0.68	0.60	
4	0.52	1.46	0.64	

	Risk_Category	Skill_1	Skill_2	Skill_3	Skill_4	Skill_5	Skill_6	\
0	High	0.45	0.10	0.46	0.33	0.14	0.65	
1	Low	0.02	0.52	0.40	0.05	0.97	0.23	
2	High	0.01	0.94	0.56	0.39	0.02	0.23	
3	Medium	0.43	0.21	0.57	0.03	0.84	0.45	
4	Medium	0.75	0.54	0.59	0.97	0.61	0.28	

	Skill_7	Skill_8	Skill_9	Skill_10
0	0.06	0.72	0.94	0.00
1	0.09	0.62	0.38	0.98
2	0.24	0.68	0.61	0.83
3	0.40	0.93	0.73	0.33
4	0.30	0.17	0.02	0.42

- Dataset Information

```
In [45]: print("\nDataset info:")
print(df.info())
```

```
Dataset info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 18 columns):
 #   Column           Non-Null Count Dtype  
 ---  ----             -----          ----- 
 0   Job_Title        3000 non-null   object  
 1   Average_Salary   3000 non-null   int64   
 2   Years_Experience 3000 non-null   int64   
 3   Education_Level 3000 non-null   object  
 4   AI_Exposure_Index 3000 non-null   float64 
 5   Tech_Growth_Factor 3000 non-null   float64 
 6   Automation_Probability_2030 3000 non-null   float64 
 7   Risk_Category    3000 non-null   object  
 8   Skill_1          3000 non-null   float64 
 9   Skill_2          3000 non-null   float64 
 10  Skill_3          3000 non-null   float64 
 11  Skill_4          3000 non-null   float64 
 12  Skill_5          3000 non-null   float64 
 13  Skill_6          3000 non-null   float64 
 14  Skill_7          3000 non-null   float64 
 15  Skill_8          3000 non-null   float64 
 16  Skill_9          3000 non-null   float64 
 17  Skill_10         3000 non-null   float64 

dtypes: float64(13), int64(2), object(3)
memory usage: 422.0+ KB
None
```

• Check Missing Values....

```
In [46]: print("\nMissing values:")
print(df.isnull().sum())
```

```
Missing values:
Job_Title          0
Average_Salary     0
Years_Experience   0
Education_Level    0
AI_Exposure_Index  0
Tech_Growth_Factor 0
Automation_Probability_2030 0
Risk_Category      0
Skill_1            0
Skill_2            0
Skill_3            0
Skill_4            0
Skill_5            0
Skill_6            0
Skill_7            0
Skill_8            0
Skill_9            0
Skill_10           0
dtype: int64
```

- **Statistical Summary mean, min, max, and standard deviation.....**

```
In [47]: print(df.describe(include="all"))
```

	Job_Title	Average_Salary	Years_Experience	Education_Level	\
count	3000	3000.00000	3000.00000	3000	
unique	20	NaN	NaN	4	
top	Software Engineer	NaN	NaN	High School	
freq	175	NaN	NaN	784	
mean	NaN	89372.279000	14.677667	NaN	
std	NaN	34608.088767	8.739788	NaN	
min	NaN	30030.000000	0.000000	NaN	
25%	NaN	58640.000000	7.000000	NaN	
50%	NaN	89318.000000	15.000000	NaN	
75%	NaN	119086.500000	22.000000	NaN	
max	NaN	149798.000000	29.000000	NaN	
	AI_Exposure_Index	Tech_Growth_Factor	Automation_Probability_2030	\	
count	3000.00000	3000.00000	3000.00000		
unique	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	
mean	0.501283	0.995343	0.501503		
std	0.284004	0.287669	0.247881		
min	0.000000	0.500000	0.050000		
25%	0.260000	0.740000	0.310000		
50%	0.500000	1.000000	0.500000		
75%	0.740000	1.240000	0.700000		
max	1.000000	1.500000	0.950000		
	Risk_Category	Skill_1	Skill_2	Skill_3	Skill_4 \
count	3000	3000.00000	3000.00000	3000.00000	3000.00000
unique	3	NaN	NaN	NaN	NaN
top	Medium	NaN	NaN	NaN	NaN
freq	1521	NaN	NaN	NaN	NaN
mean	NaN	0.496973	0.497233	0.499313	0.503667
std	NaN	0.287888	0.288085	0.288354	0.287063
min	NaN	0.000000	0.000000	0.000000	0.000000
25%	NaN	0.240000	0.250000	0.250000	0.260000
50%	NaN	0.505000	0.500000	0.500000	0.510000
75%	NaN	0.740000	0.740000	0.750000	0.750000
max	NaN	1.000000	1.000000	1.000000	1.000000
	Skill_5	Skill_6	Skill_7	Skill_8	Skill_9 \
count	3000.00000	3000.00000	3000.00000	3000.00000	3000.00000
unique	NaN	NaN	NaN	NaN	NaN
top	NaN	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN	NaN
mean	0.490270	0.499807	0.499160	0.502843	0.501433
std	0.285818	0.286050	0.288044	0.289832	0.285818
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.240000	0.260000	0.250000	0.250000	0.260000
50%	0.490000	0.500000	0.490000	0.500000	0.500000
75%	0.730000	0.740000	0.750000	0.750000	0.740000
max	1.000000	1.000000	1.000000	1.000000	1.000000
	Skill_10				
count	3000.00000				
unique	NaN				
top	NaN				

```
freq      NaN
mean    0.493627
std     0.286464
min    0.000000
25%    0.250000
50%    0.490000
75%    0.740000
max    1.000000
```

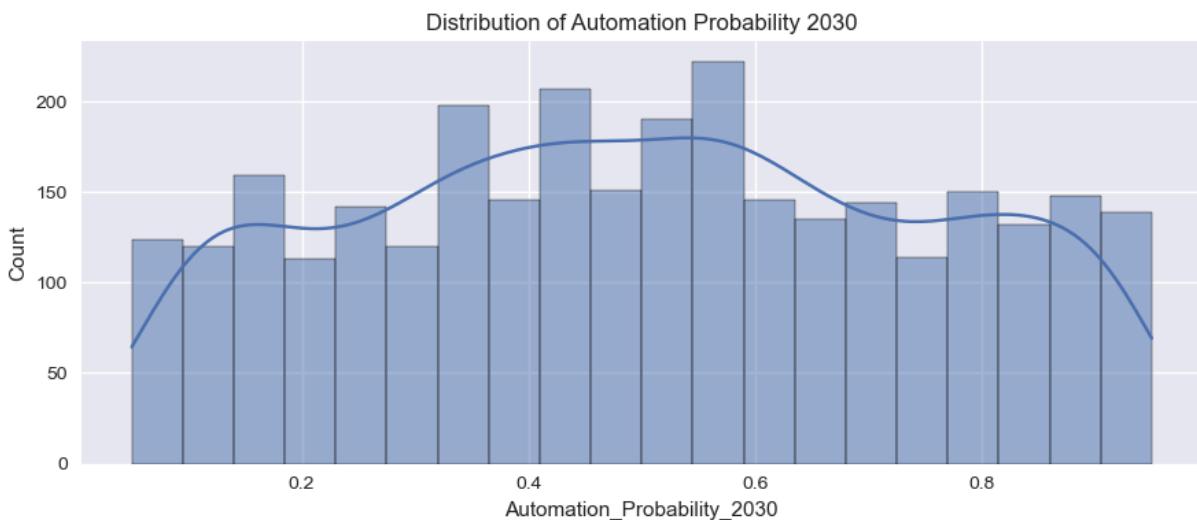
- Summary(Risk Category)

```
In [5]: print("\nSummary(Risk Category)")
print(df['Risk_Category'].value_counts())
```

```
Summary(Risk Category)
Risk_Category
Medium    1521
High      740
Low       739
Name: count, dtype: int64
```

- Histogram of Automation_Probability_2030.....

```
In [124]: plt.figure(figsize=(9, 4))
sns.histplot(df["Automation_Probability_2030"], kde=True, bins=20)
plt.title("Distribution of Automation Probability 2030")
plt.xlabel("Automation_Probability_2030")
plt.ylabel("Count")
plt.tight_layout()
plt.show()
```



- Risk category counts visualization

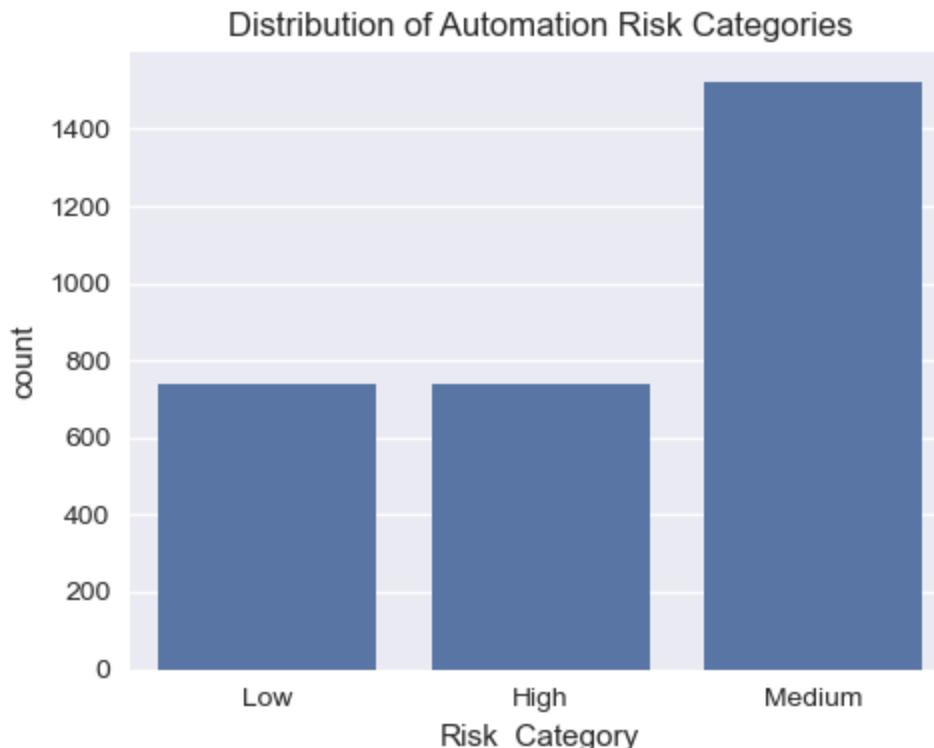
```
In [125]: print("\nRisk category counts:")
print(df['Risk_Category'].value_counts())
```

```
plt.figure(figsize=(5,4))
sns.countplot(x="Risk_Category", data=df, order=["Low", "High", "Medium"])
plt.title("Distribution of Automation Risk Categories")
plt.tight_layout()
plt.show()
```

Risk category counts:

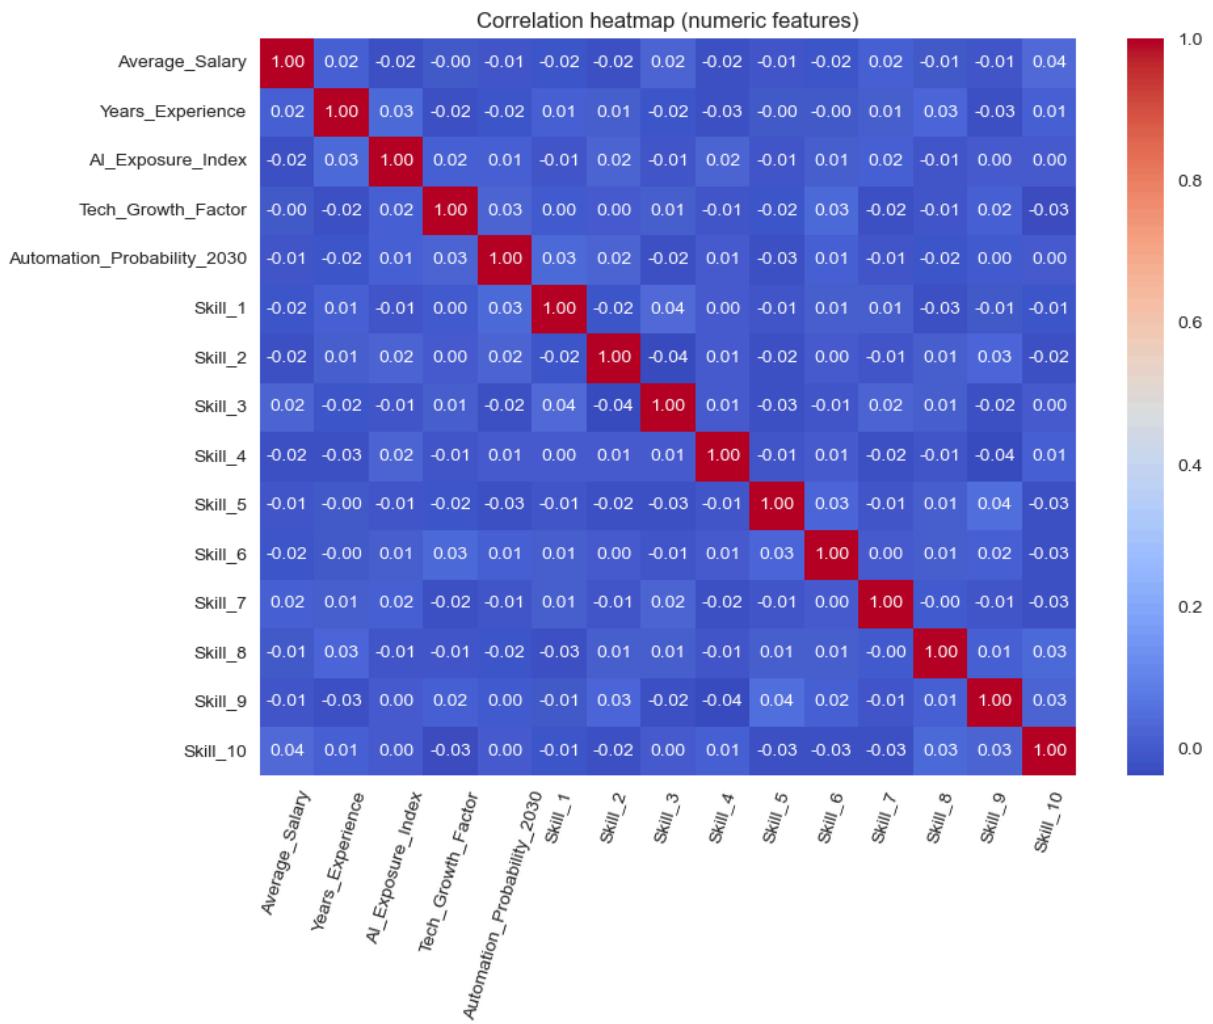
Risk Category	Count
Medium	1521
High	740
Low	739

Name: count, dtype: int64



- Correlation heatmap for numeric features.....

```
In [57]: numeric_cols = df.select_dtypes(include=[np.number]).columns
plt.figure(figsize=(10, 8))
corr = df[numeric_cols].corr()
sns.heatmap(corr, annot=True, fmt='.2f', cmap="coolwarm")
plt.title("Correlation heatmap (numeric features)")
plt.xticks(rotation=73)
plt.tight_layout()
plt.show()
```



- Encode the target variable (Risk_Category) into numerical labels.....

```
In [8]: label_encoder = LabelEncoder()
df["Risk_Category_Encoded"] = label_encoder.fit_transform(df["Risk_Category"])
```

- Separates input features (X) and the target variable (y).....

```
In [9]: feature_cols = [c for c in df.select_dtypes(include=[np.number]).columns
                   if c != "Risk_Category_Encoded"]

X = df[feature_cols]
y = df["Risk_Category_Encoded"]
```

- Split data into training 80% and testing 20% sets.....

```
In [10]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y)
```

- Training testing Shape.....

```
In [17]: print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y value counts:\n", y.value_counts())
```

```
X_train shape: (2400, 15)
X_test shape: (600, 15)
y value counts:
Risk_Category_Encoded
2      1521
0       740
1       739
Name: count, dtype: int64
```

- Standardize all numerical features using StandardScaler.....

```
In [13]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

model_results = []
trained_models = []
```

- Starting Unsupervised Analysis: K-Means & PCA....
- Elbow Method for Optimal Clusters....

```
In [68]: print("--- Starting Unsupervised Analysis: K-Means & PCA ---")
inertia_values = []
max_k = 10
# Test K from 1 to max_k
for k in range(1, max_k + 1):
    # Use X_train_scaled as K-Means is distance-based
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(X_train_scaled)
    inertia_values.append(kmeans.inertia_)
```

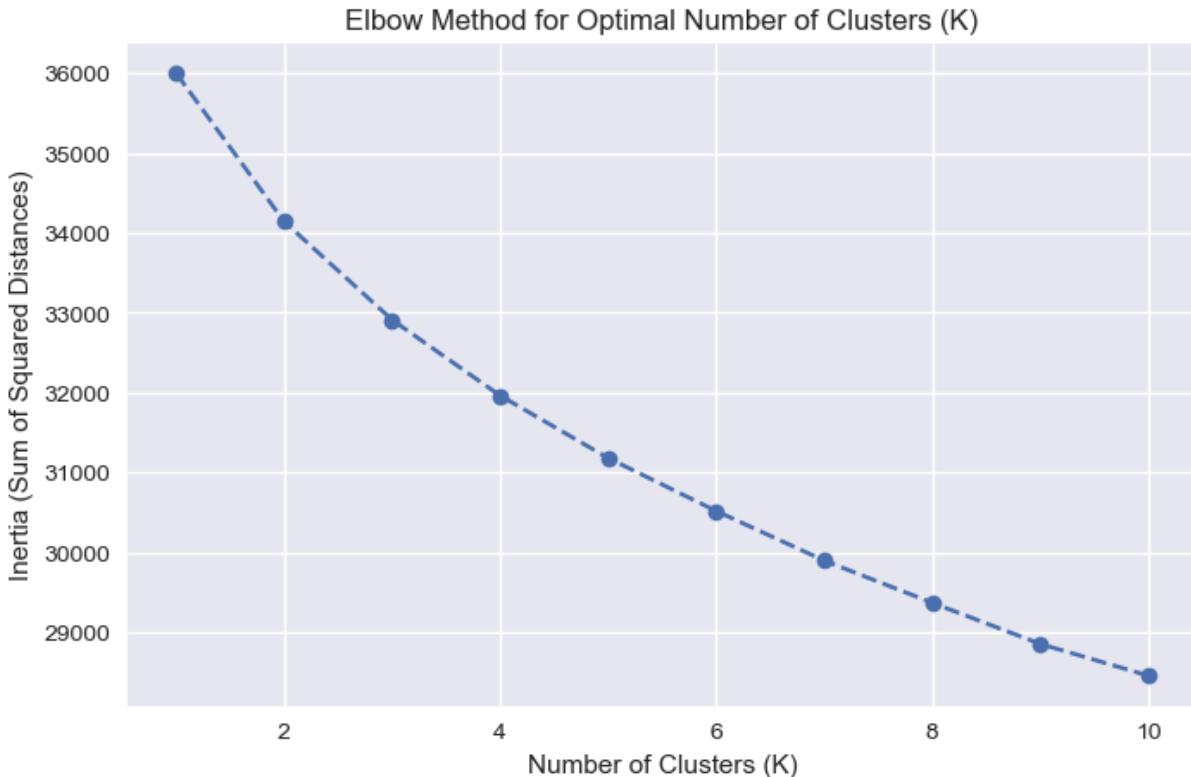
--- Starting Unsupervised Analysis: K-Means & PCA ---

1. Running Elbow Method to find optimal K...

Visualize the Elbow...

In [69]:

```
# Visualize the Elbow Method
plt.figure(figsize=(8, 5))
plt.plot(range(1, max_k + 1), inertia_values, marker='o', linestyle='--')
plt.title('Elbow Method for Optimal Number of Clusters (K)')
plt.xlabel('Number of Clusters (K)')
plt.ylabel('Inertia (Sum of Squared Distances)')
plt.grid(True)
plt.show()
print("")
```



- Clustering, PCA & Validation...

In [88]:

```
K_VAL = 3
kmeans_val = KMeans(n_clusters=K_VAL, random_state=42, n_init='auto')
cluster_labels = kmeans_val.fit_predict(X_train_scaled)

# Validation: Compare Clusters to True Risk Categories
comparison_df = pd.DataFrame({
    'True_Risk_Category_Encoded': y_train,
    'KMeans_Cluster': cluster_labels
})
cross_tab = pd.crosstab(comparison_df['True_Risk_Category_Encoded'], comparison_df['KMeans_Cluster'])

print(f"Cluster Alignment (K={K_VAL}):")
print(cross_tab)
print("\n")
```

--- Unsupervised Analysis: Step 2 (Clustering, 2D PCA & Validation) ---

Cluster Alignment (K=3):

KMeans_Cluster	0	1	2
True_Risk_Category_Encoded			
0	104	260	228
1	323	113	155
2	412	380	425

- Perform Principal Component Analysis (PCA) to reduce dimensions...

```
In [91]: # 2. Dimensionality Reduction with PCA
pca = PCA(n_components=3, random_state=42)
X_pca = pca.fit_transform(X_train_scaled)

explained_variance = pca.explained_variance_ratio_.sum()
print(f"Variance Explained by 3 PCs: {explained_variance:.2%}")
```

Variance Explained by 3 PCs: 22.32%

- Visualization for K-Means clusters in a 3D scatter plot...

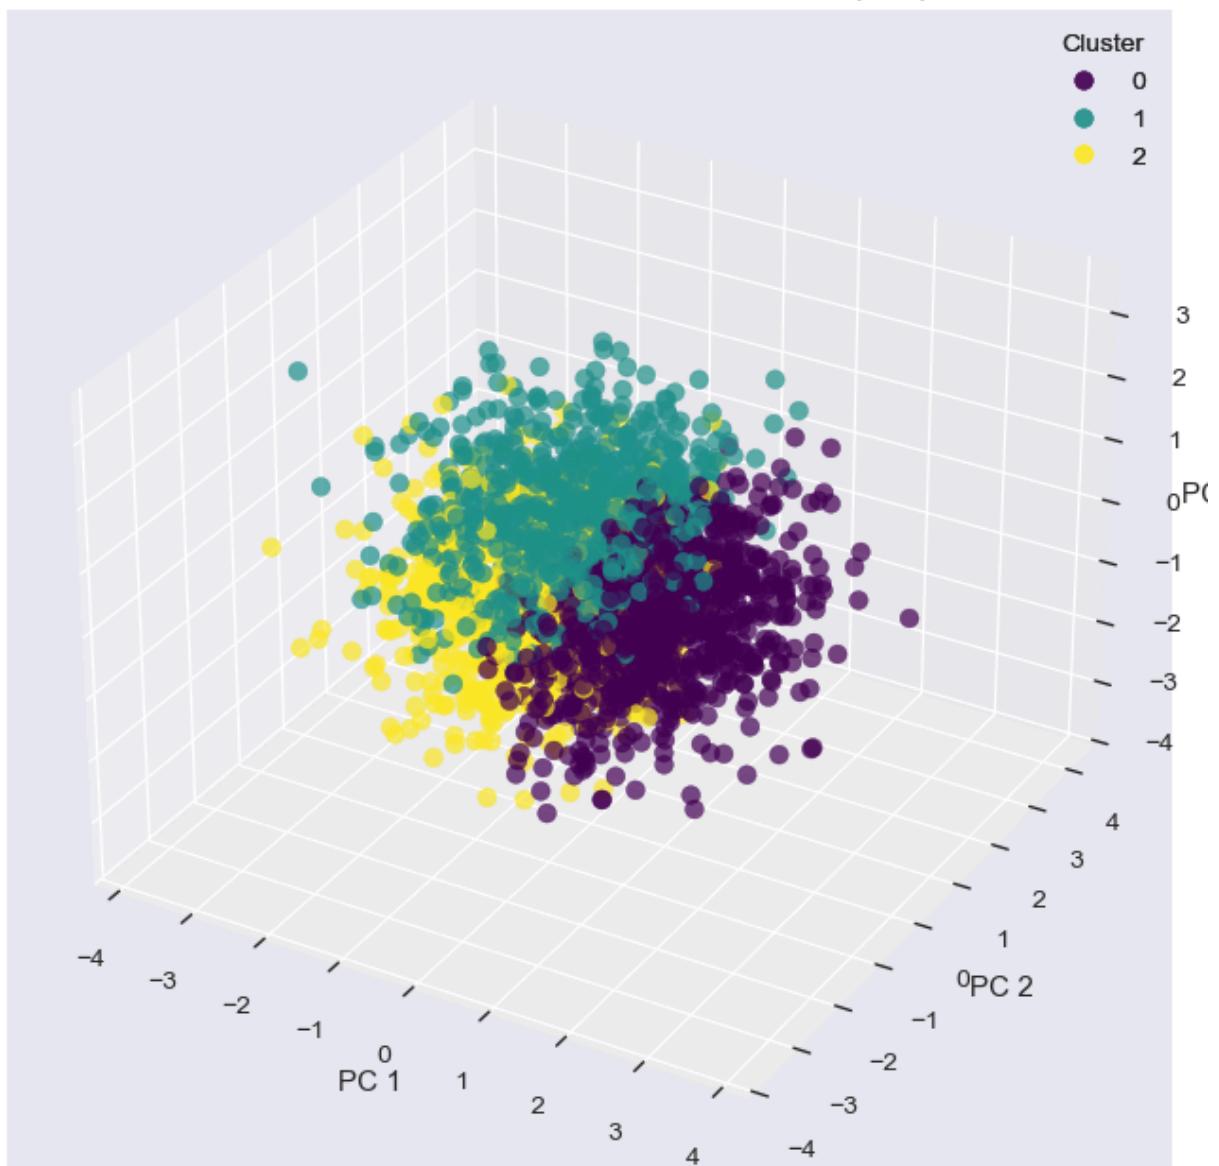
```
In [92]: # 3. Create 3D Visualization
pca_df = pd.DataFrame(data=X_pca, columns=['PC1', 'PC2', 'PC3'])
pca_df['KMeans_Cluster'] = cluster_labels
pca_df['KMeans_Cluster'] = pca_df['KMeans_Cluster'].astype('category')

fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')

scatter = ax.scatter(pca_df['PC1'], pca_df['PC2'], pca_df['PC3'],
                     c=pca_df['KMeans_Cluster'].cat.codes,
                     cmap='viridis',
                     s=50,
                     alpha=0.7)

ax.set_title(f'3D PCA Visualization of K-Means Clusters (K={K_VAL})')
ax.set_xlabel('PC 1')
ax.set_ylabel('PC 2')
ax.set_zlabel('PC 3')
legend1 = ax.legend(*scatter.legend_elements(), title="Cluster")
ax.add_artist(legend1)
plt.show()
print("")
```

3D PCA Visualization of K-Means Clusters (K=3)



- **supervise Learning: Define a reusable function to train, predict, and evaluate model performance.....**

```
In [41]: def evaluate_model(name, model, X_tr, y_tr, X_te, y_te):
    model.fit(X_tr, y_tr)
    y_pred = model.predict(X_te)
    acc = accuracy_score(y_te, y_pred)
    model_results[name] = acc
    trained_models[name] = model
    print("= " + name + "=")
    print("Accuracy:", round(acc, 4))
    print(classification_report(y_te, y_pred, target_names=label_encoder.classes_))

    # Confusion matrix
    cm = confusion_matrix(y_test, y_pred)
```

```

plt.figure(figsize=(5,4))
sns.heatmap(
    cm,
    annot=True,
    fmt="d",
    cmap="Blues",
    xticklabels=label_encoder.classes_,
    yticklabels=label_encoder.classes_
)
plt.title("Confusion Matrix - " + name)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.tight_layout()
plt.show()
print()
return model, acc

```

- Logistic Regression Predictions
- Logistic Regression Accuracy And
- Confusion Matrix, f1-score for Logistic Regression
- ...

In [67]:

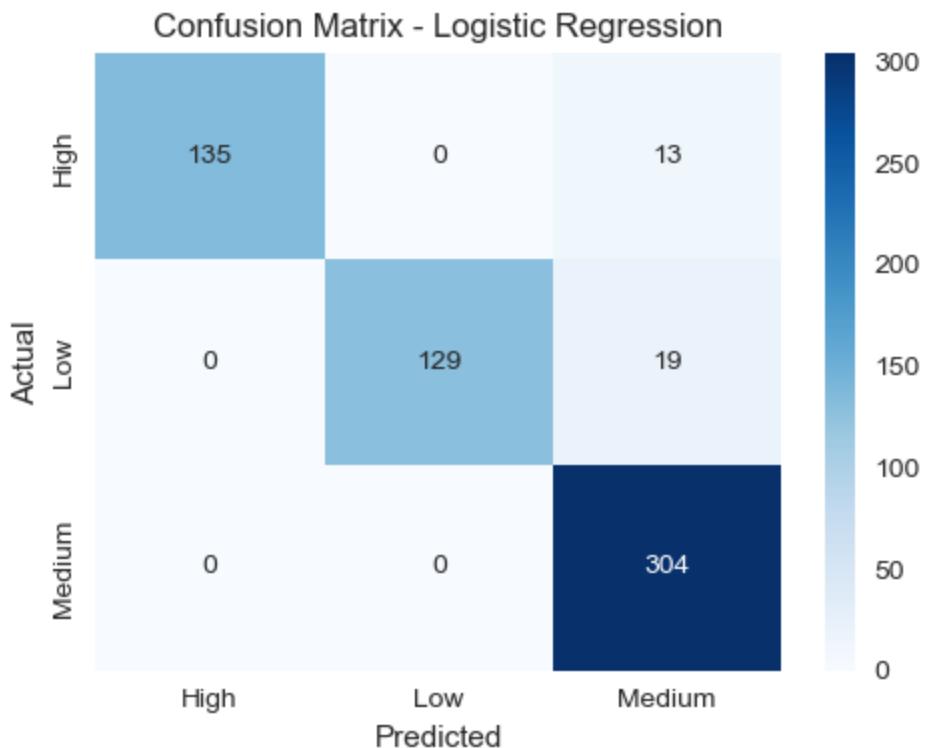
```

# Logistic Regression
log_reg = LogisticRegression(max_iter=1000, C=0.0084,)
log_reg, log_acc = evaluate_model(
    "Logistic Regression", log_reg, X_train_scaled, y_train, X_test_scaled, y_test
)
print("Final Logistic Regression accuracy:", log_acc)

=Logistic Regression=
Accuracy: 0.9467

```

	precision	recall	f1-score	support
High	1.00	0.91	0.95	148
Low	1.00	0.87	0.93	148
Medium	0.90	1.00	0.95	304
accuracy			0.95	600
macro avg	0.97	0.93	0.95	600
weighted avg	0.95	0.95	0.95	600



Final Logistic Regression accuracy: 0.9466666666666667

- KNN Predictions
- KNN Accuracy
- Confusion Matrix, f1-score for KNN

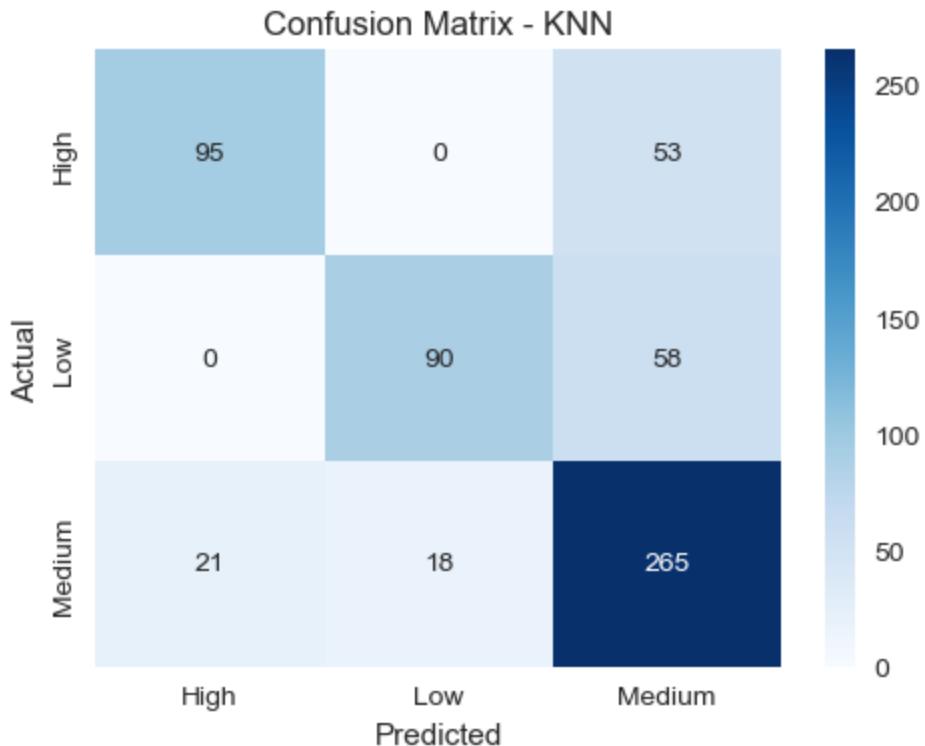
In [43]:

```
# KNN
knn = KNeighborsClassifier(n_neighbors=7)
knn, knn_acc = evaluate_model(
    "KNN", knn, X_train_scaled, y_train, X_test_scaled, y_test
)
```

=KNN=

Accuracy: 0.75

	precision	recall	f1-score	support
High	0.82	0.64	0.72	148
Low	0.83	0.61	0.70	148
Medium	0.70	0.87	0.78	304
accuracy			0.75	600
macro avg	0.79	0.71	0.73	600
weighted avg	0.76	0.75	0.75	600



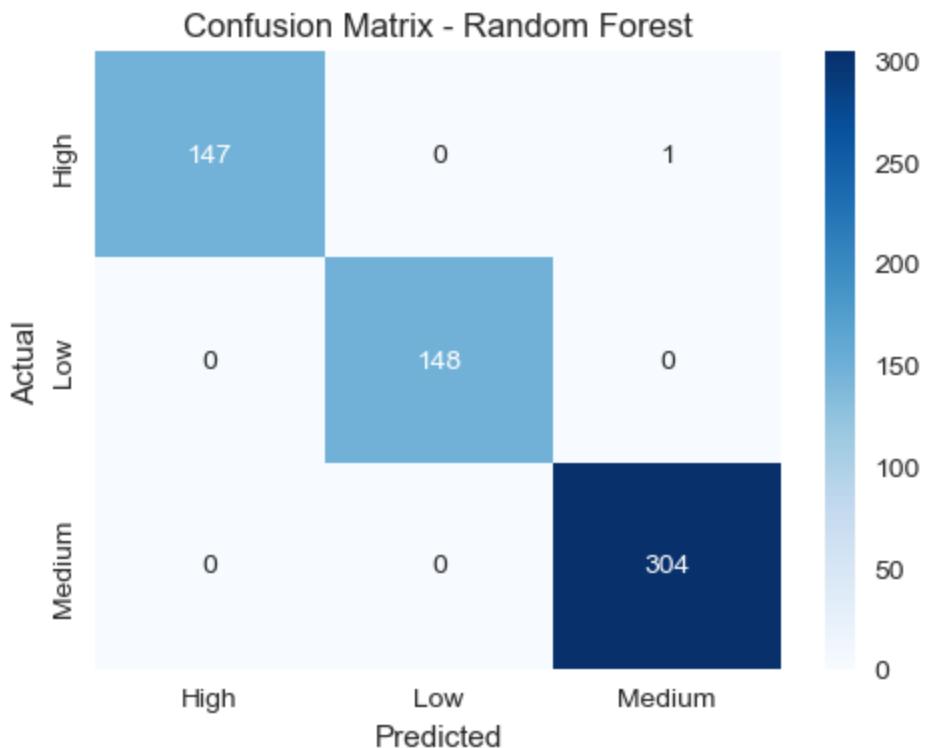
- Random Forest Predictions..
- Random Forest Accuracy..
- Confusion Matrix, f1-score for Random Forest..

```
In [75]: # Random Forest (no scaling needed)
rf = RandomForestClassifier(
    n_estimators=150, max_depth=3, random_state=42
)
rf, rf_acc = evaluate_model(
    "Random Forest", rf, X_train, y_train, X_test, y_test
)
```

```
=Random Forest=
Accuracy: 0.9983
      precision    recall  f1-score   support

      High       1.00     0.99       1.00      148
      Low        1.00     1.00       1.00      148
   Medium       1.00     1.00       1.00      304

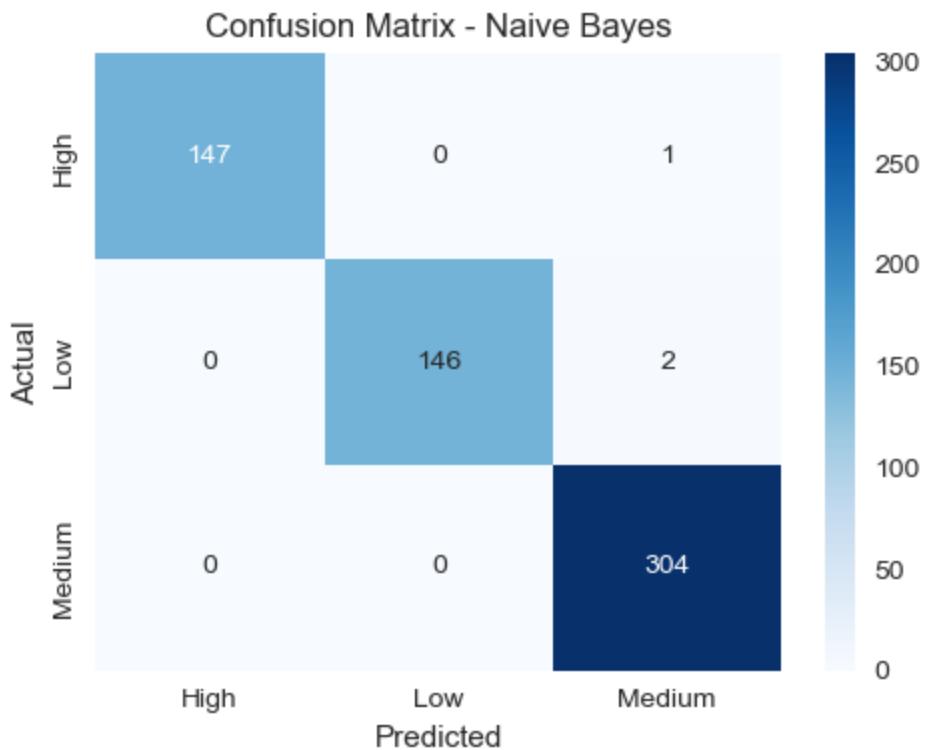
  accuracy                           1.00      600
  macro avg       1.00     1.00       1.00      600
weighted avg       1.00     1.00       1.00      600
```



- Naive Bayes Model Training
- Naive Bayes Predictions
- Naive Bayes Accuracy, confusion matrix and f1-score

```
In [45]: # Naive Bayes
nb = GaussianNB()
nb, nb_acc = evaluate_model(
    "Naive Bayes", nb, X_train_scaled, y_train, X_test_scaled, y_test
)
```

```
=Naive Bayes=
Accuracy: 0.995
      precision    recall  f1-score   support
High       1.00     0.99     1.00     148
Low       1.00     0.99     0.99     148
Medium     0.99     1.00     1.00     304
accuracy          0.99     0.99     0.99     600
macro avg       1.00     0.99     0.99     600
weighted avg     1.00     0.99     0.99     600
```



- Support Vector Machine (SVM) Model Training...
- Support Vector Machine (SVM) Predictions....
- Support Vector Machine (SVM) Bayes Accuracy...

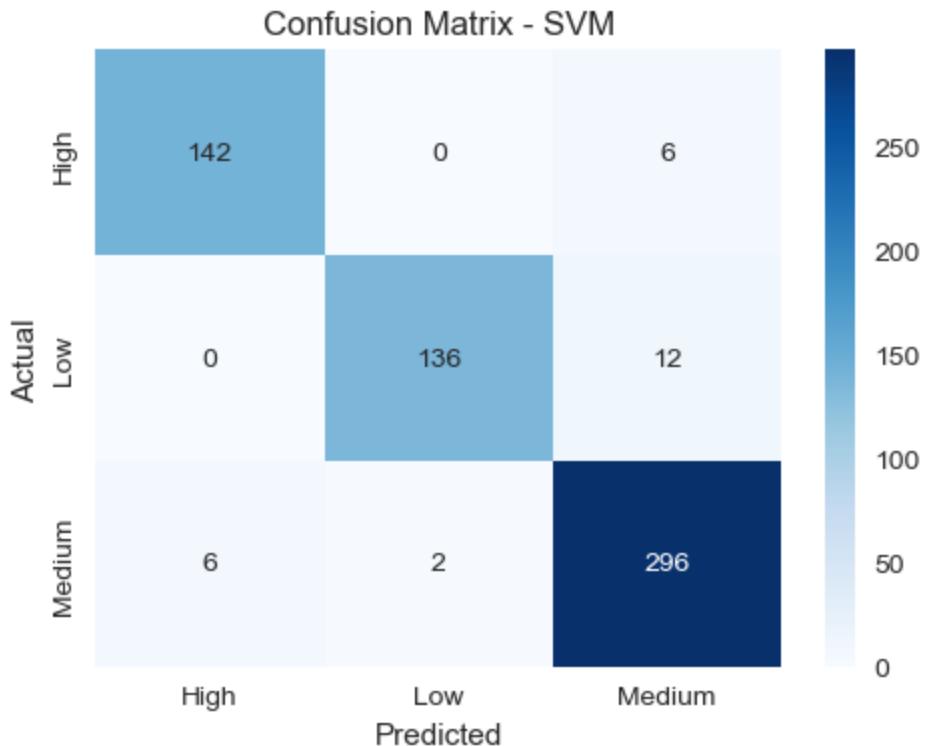
In [46]:

```
# SVM
svm_clf = SVC(kernel="rbf", C=1.0, gamma="scale", probability=True, random_state=42)
svm_clf, svm_acc = evaluate_model(
    "SVM", svm_clf, X_train_scaled, y_train, X_test_scaled, y_test
)
```

=SVM=

Accuracy: 0.9567

	precision	recall	f1-score	support
High	0.96	0.96	0.96	148
Low	0.99	0.92	0.95	148
Medium	0.94	0.97	0.96	304
accuracy			0.96	600
macro avg	0.96	0.95	0.96	600
weighted avg	0.96	0.96	0.96	600

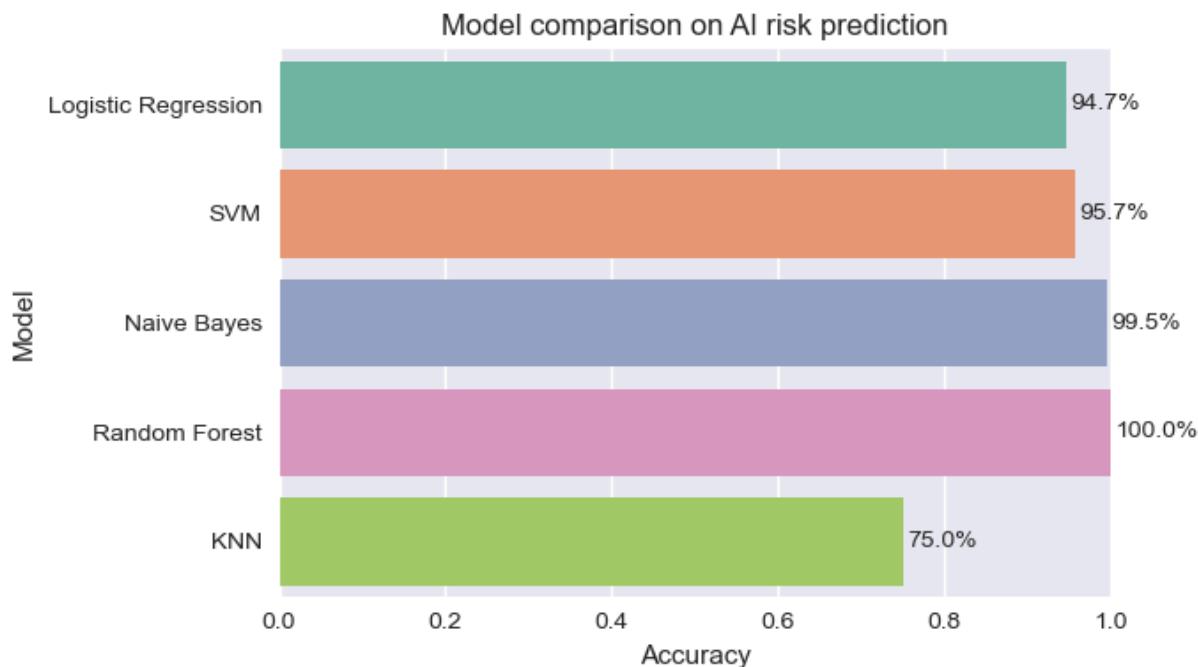


- Model Performance Summary Table visualization.....

```
In [47]: print("Model accuracies summary:")
for name, acc in model_results.items():
    print(name + ":", round(acc, 4))
```

Model accuracies summary:
 Logistic Regression: 0.9467
 SVM: 0.9567
 Naive Bayes: 0.995
 Random Forest: 1.0
 KNN: 0.75

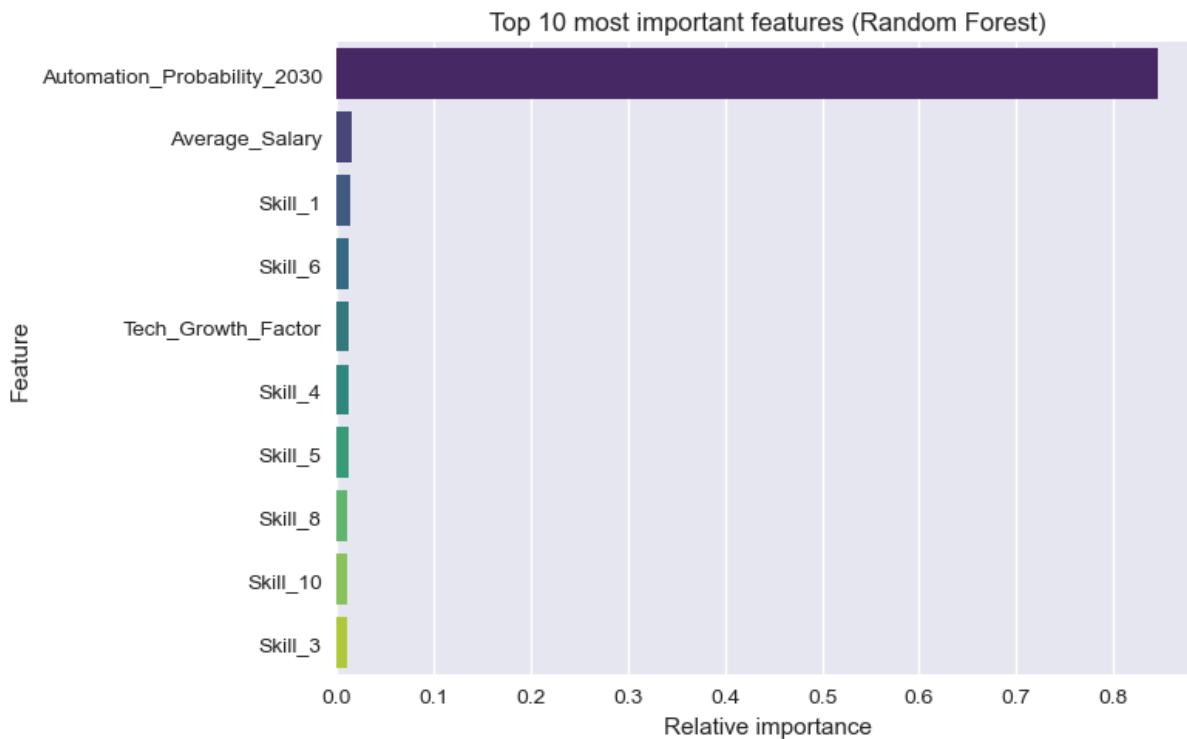
```
In [48]: plt.figure(figsize=(7, 4))
model_names = list(model_results.keys())
accuracies = [model_results[m] for m in model_names]
sns.barplot(x=accuracies, y=model_names, palette="Set2")
for i, v in enumerate(accuracies):
    plt.text(v + 0.005, i, str(round(v * 100, 1)) + "%", va="center")
plt.xlim(0, 1)
plt.xlabel("Accuracy")
plt.ylabel("Model")
plt.title("Model comparison on AI risk prediction")
plt.tight_layout()
plt.show()
```



- Feature importance from Random Forest....

```
In [69]: # Feature importance from Random Forest
importances = rf.feature_importances_
feat_names = X.columns
feat_imp = list(zip(feat_names, importances))
feat_imp_sorted = sorted(feat_imp, key=lambda x: x[1], reverse=True)[:10]
feat_labels = [f[0] for f in feat_imp_sorted]
feat_values = [f[1] for f in feat_imp_sorted]
```

```
In [27]: plt.figure(figsize=(8, 5))
sns.barplot(x=feat_values, y=feat_labels, palette="viridis")
plt.title("Top 10 most important features (Random Forest)")
plt.xlabel("Relative importance")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```



- visualize the Learning Curve for the best model (Random Forest)....

```
In [76]: best_model_name = max(model_results, key=model_results.get)
best_model = trained_models[best_model_name]
print("Best model by accuracy:", best_model_name)

# Choose correct feature representation
if best_model_name in ["Logistic Regression", "Random Forest", "KNN", "Naive Bayes",
    X_curve = X_train_scaled
else:
    X_curve = X_train

train_sizes, train_scores, test_scores = learning_curve(
    best_model,
    X_curve,
    y_train,
    cv=5,
    train_sizes=np.linspace(0.1, 1.0, 5),
    scoring="accuracy",
    n_jobs=-1
)

train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)
val_mean = np.mean(test_scores, axis=1)
val_std = np.std(test_scores, axis=1)

plt.figure(figsize=(8, 5))
plt.plot(train_sizes, train_mean, "o-", color="tab:blue", label="Training score")
plt.fill_between(train_sizes, train_mean - train_std,
```

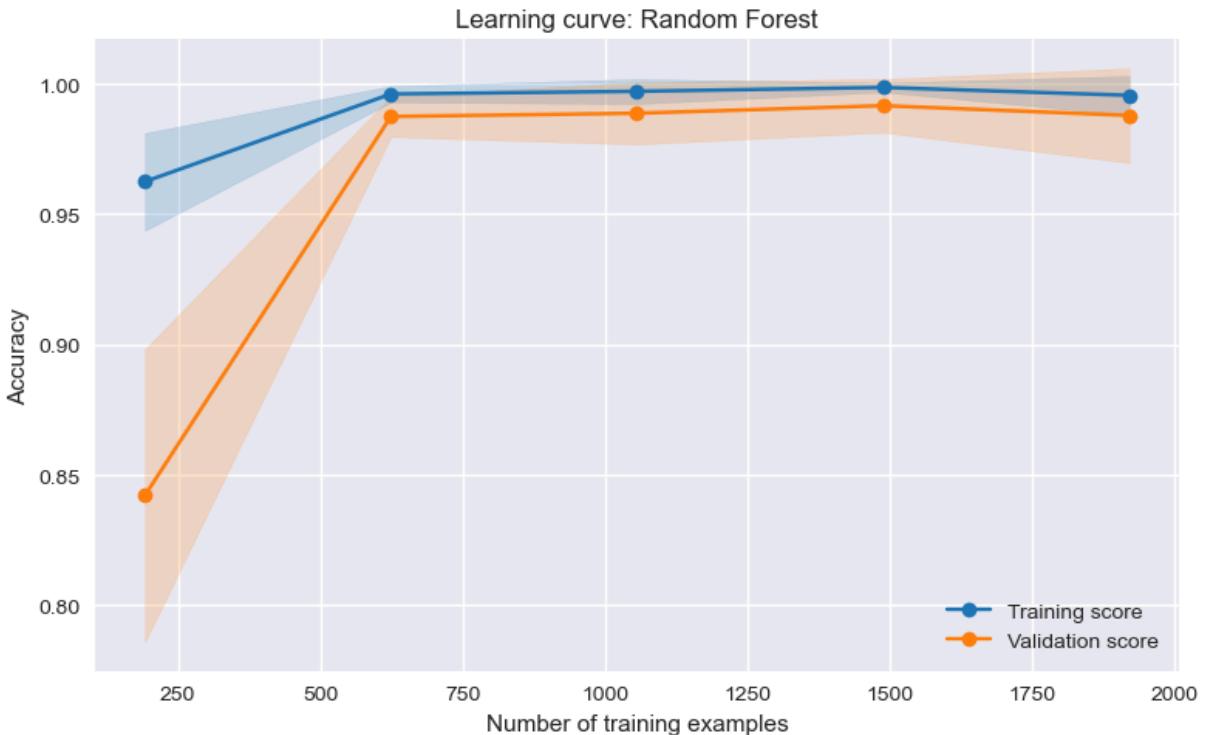
```

        train_mean + train_std, alpha=0.2, color="tab:blue")
plt.plot(train_sizes, val_mean, "o-", color="tab:orange", label="Validation score")
plt.fill_between(train_sizes, val_mean - val_std,
                 val_mean + val_std, alpha=0.2, color="tab:orange")

plt.title("Learning curve: " + best_model_name)
plt.xlabel("Number of training examples")
plt.ylabel("Accuracy")
plt.legend(loc="best")
plt.tight_layout()
plt.show()

```

Best model by accuracy: Random Forest



- model testing

```
In [33]: # Create a dictionary for a new job profile (e.g., "Machine Operator")
new_job_data = {
    'Average_Salary': [45000],
    'Years_Experience': [10],
    'AI_Exposure_Index': [0.40],
    'Tech_Growth_Factor': [0.15],
    'Automation_Probability_2030': [0.92], # Critical for High Risk
    'Skill_1': [0.15], # Low for Cognitive Skills
    'Skill_2': [0.25],
    'Skill_3': [0.85], # High for Manual/Repetitive Skills
    'Skill_4': [0.20],
    'Skill_5': [0.70],
    'Skill_6': [0.35],
    'Skill_7': [0.20],
    'Skill_8': [0.10],
    'Skill_9': [0.50],
}
```

```
'Skill_10': [0.90]
}
new_job_df = pd.DataFrame(new_job_data)
```

```
In [36]: new_job_scaled = scaler.transform(new_job_df)
prediction_encoded = best_model.predict(new_job_scaled)
predicted_risk = label_encoder.inverse_transform(prediction_encoded)

print("Input Automation Probability: {new_job_data['Automation_Probability_2030'][0]}
print("The Predicted AI Risk Category for the new job is: {predicted_risk[0]} ")
```

```
Input Automation Probability: 0.92
The Predicted AI Risk Category for the new job is: High
```

```
In [ ]:
```